

Import Modules

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('Black Friday Sales.csv')
```

```
In [3]: # Descriptive analysis
```

```
In [4]: df.head()
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Prod
0	1000001	P00069042	F	0-17	10	A	2	0	
1	1000001	P00248942	F	0-17	10	A	2	0	
2	1000001	P00087842	F	0-17	10	A	2	0	
3	1000001	P00085442	F	0-17	10	A	2	0	
4	1000002	P00285442	M	55+	16	C	4+	0	

```
In [5]: # Datatype info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                      550068 non-null  int64
8   Product_Category_1                  550068 non-null  int64
9   Product_Category_2                  376430 non-null  float64
10  Product_Category_3                  166821 non-null  float64
11  Purchase                            550068 non-null  int64
dtypes: float64(2), int64(5), object(5)
memory usage: 50.4+ MB
```

```
In [6]: # statistical info
df.describe()
```

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3
--	---------	------------	----------------	--------------------	--------------------	--------------------

count	5.500680e+05	550068.000000	550068.000000	550068.000000	376430.000000	166821.000000
mean	1.003029e+06	8.076707	0.409653	5.404270	9.842329	12.668243
std	1.727592e+03	6.522660	0.491770	3.936211	5.086590	4.125338
min	1.000001e+06	0.000000	0.000000	1.000000	2.000000	3.000000
25%	1.001516e+06	2.000000	0.000000	1.000000	5.000000	9.000000
50%	1.003077e+06	7.000000	0.000000	5.000000	9.000000	14.000000
75%	1.004478e+06	14.000000	1.000000	8.000000	15.000000	16.000000
max	1.006040e+06	20.000000	1.000000	20.000000	18.000000	18.000000

```
In [7]: # Unique values
df.apply(lambda x: len(x.unique()))
```


```
Out[7]: User_ID          5891
Product_ID       3631
Gender           2
Age              7
Occupation       21
City_Category    3
Stay_In_Current_City_Years  5
Marital_Status   2
Product_Category_1  20
Product_Category_2  18
Product_Category_3  16
Purchase        18105
dtype: int64
```

```
In [8]: # Null values
df.isnull().sum()
```


```
Out[8]: User_ID          0
Product_ID          0
Gender              0
Age                 0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1  0
Product_Category_2  173638
Product_Category_3  383247
Purchase            0
dtype: int64
```

Exploratory Data Ananalysis

```
In [9]: plt.figure(figsize=(10,6))
sns.distplot(df['Purchase'],bins=25);
```




```
In [10]: sns.boxplot(df['Purchase']);
```




```
In [11]: # We can see outliers in Purchase column
```


```
In [12]: # dist of numeric variablres
sns.countplot(df.Gender);
```




```
In [13]: sns.countplot(df.Age);
```




```
In [14]: sns.countplot(df.Marital_Status);
```




```
In [15]: plt.figure(figsize=(10,6))  
sns.countplot(df.Occupation);
```




```
In [16]: plt.figure(figsize=(10,6))  
sns.countplot(df.Product_Category_1);
```




```
In [17]: plt.figure(figsize=(10,6))  
sns.countplot(df.Product_Category_2);
```



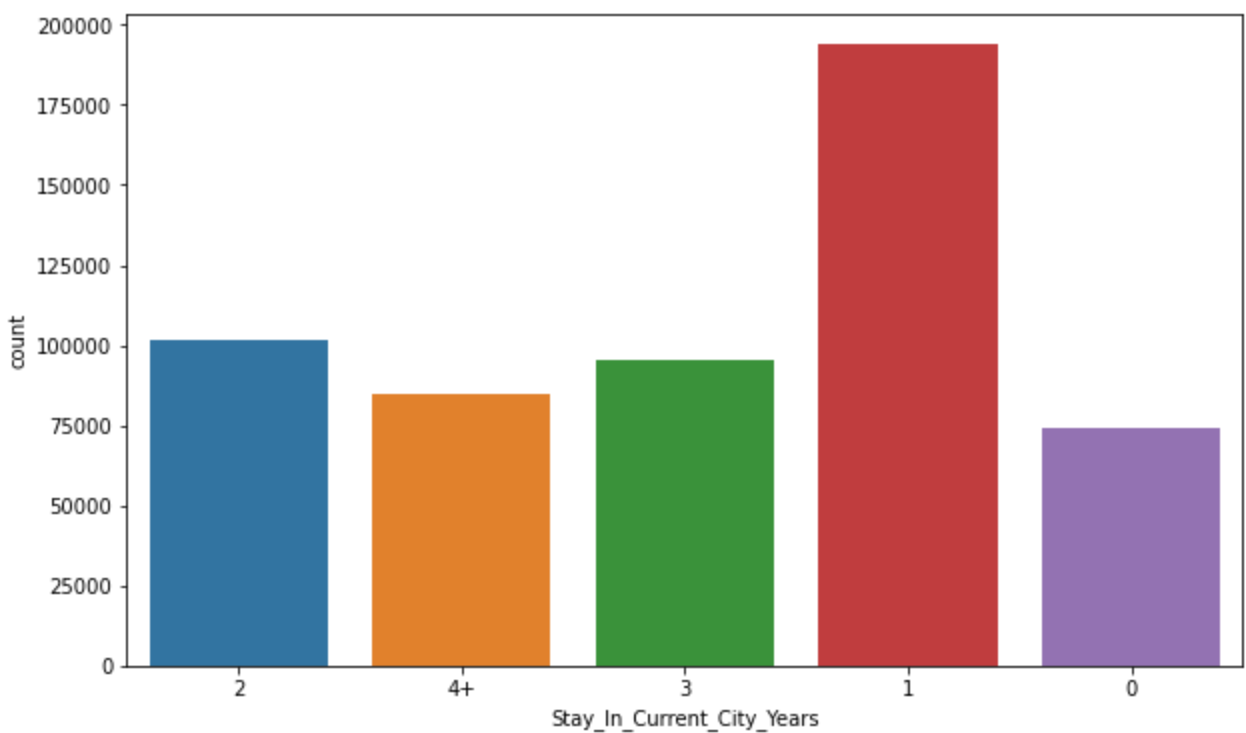
```
In [18]: plt.figure(figsize=(10,6))  
sns.countplot(df.Product_Category_3);
```




```
In [19]: plt.figure(figsize=(10,6))
sns.countplot(df.City_Category);
```




```
In [20]: plt.figure(figsize=(10,6))
sns.countplot(df.Stay_In_Current_City_Years);
```



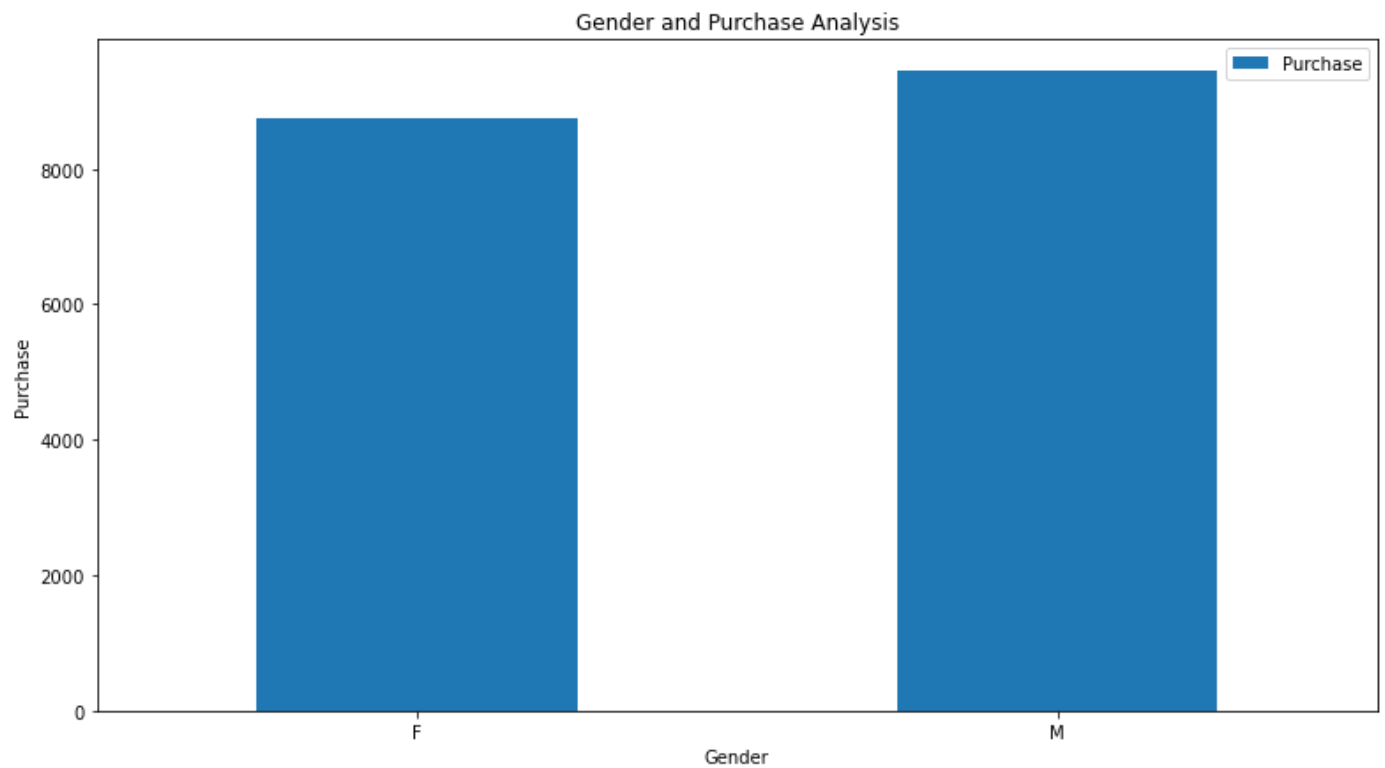
```
In [21]: # Bivariate analysis
occupation_plot = df.pivot_table(index='Occupation', values='Purchase', aggfunc=np.mean)
occupation_plot.plot(kind='bar', figsize=(10, 6))
plt.xlabel('Occupation')
plt.ylabel("Purchase")
plt.title("Occupation and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [22]: age_plot = df.pivot_table(index='Age', values='Purchase', aggfunc=np.mean)
age_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('Age')
plt.ylabel("Purchase")
plt.title("Age and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [23]: gender_plot = df.pivot_table(index='Gender', values='Purchase', aggfunc=np.mean)
gender_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('Gender')
plt.ylabel("Purchase")
plt.title("Gender and Purchase Analysis")
plt.xticks(rotation=0)
plt.show()
```



```
In [24]: df.head()
```

```
Out[24]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Prod
0	1000001	P00069042	F	0-17	10	A	2	0	

1	1000001	P00248942	F	0-17	10	A	2	0
2	1000001	P00087842	F	0-17	10	A	2	0
3	1000001	P00085442	F	0-17	10	A	2	0
4	1000002	P00285442	M	55+	16	C	4+	0

Preprocessing the Dataset

In [25]: `df.isnull().sum()`

Out[25]:

```
User_ID          0
Product_ID       0
Gender           0
Age             0
Occupation       0
City_Category    0
Stay_In_Current_City_Years  0
Marital_Status   0
Product_Category_1  0
Product_Category_2 173638
Product_Category_3 383247
Purchase         0
dtype: int64
```

In [26]: `for i in df.columns:`
`print(i," = ",len(df[i].unique()))`

```
User_ID   = 5891
Product_ID = 3631
Gender    = 2
Age       = 7
Occupation = 21
City_Category = 3
Stay_In_Current_City_Years = 5
Marital_Status = 2
Product_Category_1 = 20
Product_Category_2 = 18
Product_Category_3 = 16
Purchase   = 18105
```

In [27]: `# Remove outliers using IQR technique`

In [28]: `cols = ['Purchase']`


```
Q1 = df[cols].quantile(0.25)
Q3 = df[cols].quantile(0.75)
IQR = Q3 - Q1

df = df[~((df[cols] < (Q1 - 1.5 * IQR)) | (df[cols] > (Q3 + 1.5 * IQR))).any(axis=1)]
```

In []:

In []:

In [29]: `sns.boxplot(df['Purchase']);`




```
In [30]: # Dropping columns
```

```
In [31]: df['Product_Category_2'] = df['Product_Category_2'].fillna((df['Product_Category_2']).me
df.drop(['Product_Category_3'], axis=1, inplace=True)
```

Corealtion Matrix

```
In [32]: corr = df.corr()
plt.figure(figsize=(13,6))
sns.heatmap(corr, annot=True);
```



```
In [33]: # df_Gender = pd.get_dummies(train['Gender'])
# df_Age = pd.get_dummies(train['Age'])
# df_City_Category = pd.get_dummies(train['City_Category'])
# df_Stay_In_Current_City_Years = pd.get_dummies(train['Stay_In_Current_City_Years'])

# data_final= pd.concat([train, df_Gender, df_Age, df_City_Category, df_Stay_In_Current_
# data_final.head()
```

```
In [34]: # from sklearn.preprocessing import LabelEncoder
# LE= LabelEncoder()

In [35]: # df['Gender'] = LE.fit_transform(df['Gender'])
# df['Age'] = LE.fit_transform(df['Age'])
# df['City_Category'] = LE.fit_transform(df['City_Category'])
# df['Stay_In_Current_City_Years'] = LE.fit_transform(df['Stay_In_Current_City_Years'])

In [36]: df= pd.get_dummies(df, columns = ['Gender', 'Age', 'City_Category', 'Stay_In_Current_City_

In [37]: df.head()
```

Out[37]:

	User_ID	Product_ID	Marital_Status	Purchase	Gender_F	Gender_M	Age_0-17	Age_18-25	Age_26-35	Age_36-45	...
0	1000001	P00069042	0	8370	1	0	1	0	0	0	...
1	1000001	P00248942	0	15200	1	0	1	0	0	0	...
2	1000001	P00087842	0	1422	1	0	1	0	0	0	...
3	1000001	P00085442	0	1057	1	0	1	0	0	0	...
4	1000002	P00285442	0	7969	0	1	0	0	0	0	...

5 rows × 79 columns

Input Split

```
In [38]: df.head()
```

Out[38]:

	User_ID	Product_ID	Marital_Status	Purchase	Gender_F	Gender_M	Age_0-17	Age_18-25	Age_26-35	Age_36-45	...
0	1000001	P00069042	0	8370	1	0	1	0	0	0	...
1	1000001	P00248942	0	15200	1	0	1	0	0	0	...
2	1000001	P00087842	0	1422	1	0	1	0	0	0	...
3	1000001	P00085442	0	1057	1	0	1	0	0	0	...
4	1000002	P00285442	0	7969	0	1	0	0	0	0	...

5 rows × 79 columns

```
In [39]: X = df.drop(columns=['User_ID', 'Product_ID', 'Purchase'])
y = df['Purchase']
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 547391 entries, 0 to 550067
Data columns (total 76 columns):
#      Column                                Non-Null Count  Dtype
---  -
0      Marital_Status                547391 non-null  int64
1      Gender_F                      547391 non-null  uint8
2      Gender_M                      547391 non-null  uint8
3      Age_0-17                     547391 non-null  uint8
4      Age_18-25                     547391 non-null  uint8
5      Age_26-35                     547391 non-null  uint8
```

6	Age_36-45	547391	non-null	uint8
7	Age_46-50	547391	non-null	uint8
8	Age_51-55	547391	non-null	uint8
9	Age_55+	547391	non-null	uint8
10	City_Category_A	547391	non-null	uint8
11	City_Category_B	547391	non-null	uint8
12	City_Category_C	547391	non-null	uint8
13	Stay_In_Current_City_Years_0	547391	non-null	uint8
14	Stay_In_Current_City_Years_1	547391	non-null	uint8
15	Stay_In_Current_City_Years_2	547391	non-null	uint8
16	Stay_In_Current_City_Years_3	547391	non-null	uint8
17	Stay_In_Current_City_Years_4+	547391	non-null	uint8
18	Occupation_0	547391	non-null	uint8
19	Occupation_1	547391	non-null	uint8
20	Occupation_2	547391	non-null	uint8
21	Occupation_3	547391	non-null	uint8
22	Occupation_4	547391	non-null	uint8
23	Occupation_5	547391	non-null	uint8
24	Occupation_6	547391	non-null	uint8
25	Occupation_7	547391	non-null	uint8
26	Occupation_8	547391	non-null	uint8
27	Occupation_9	547391	non-null	uint8
28	Occupation_10	547391	non-null	uint8
29	Occupation_11	547391	non-null	uint8
30	Occupation_12	547391	non-null	uint8
31	Occupation_13	547391	non-null	uint8
32	Occupation_14	547391	non-null	uint8
33	Occupation_15	547391	non-null	uint8
34	Occupation_16	547391	non-null	uint8
35	Occupation_17	547391	non-null	uint8
36	Occupation_18	547391	non-null	uint8
37	Occupation_19	547391	non-null	uint8
38	Occupation_20	547391	non-null	uint8
39	Product_Category_1_1	547391	non-null	uint8
40	Product_Category_1_2	547391	non-null	uint8
41	Product_Category_1_3	547391	non-null	uint8
42	Product_Category_1_4	547391	non-null	uint8
43	Product_Category_1_5	547391	non-null	uint8
44	Product_Category_1_6	547391	non-null	uint8
45	Product_Category_1_7	547391	non-null	uint8
46	Product_Category_1_8	547391	non-null	uint8
47	Product_Category_1_9	547391	non-null	uint8
48	Product_Category_1_10	547391	non-null	uint8
49	Product_Category_1_11	547391	non-null	uint8
50	Product_Category_1_12	547391	non-null	uint8
51	Product_Category_1_13	547391	non-null	uint8
52	Product_Category_1_14	547391	non-null	uint8
53	Product_Category_1_15	547391	non-null	uint8
54	Product_Category_1_16	547391	non-null	uint8
55	Product_Category_1_17	547391	non-null	uint8
56	Product_Category_1_18	547391	non-null	uint8
57	Product_Category_1_19	547391	non-null	uint8
58	Product_Category_1_20	547391	non-null	uint8
59	Product_Category_2_2	547391	non-null	uint8
60	Product_Category_2_3	547391	non-null	uint8
61	Product_Category_2_4	547391	non-null	uint8
62	Product_Category_2_5	547391	non-null	uint8
63	Product_Category_2_6	547391	non-null	uint8
64	Product_Category_2_7	547391	non-null	uint8
65	Product_Category_2_8	547391	non-null	uint8
66	Product_Category_2_9	547391	non-null	uint8
67	Product_Category_2_10	547391	non-null	uint8
68	Product_Category_2_11	547391	non-null	uint8
69	Product_Category_2_12	547391	non-null	uint8
70	Product_Category_2_13	547391	non-null	uint8
71	Product_Category_2_14	547391	non-null	uint8

```

72 Product_Category_2_15      547391 non-null  uint8
73 Product_Category_2_16      547391 non-null  uint8
74 Product_Category_2_17      547391 non-null  uint8
75 Product_Category_2_18      547391 non-null  uint8
dtypes: int64(1), uint8(75)
memory usage: 47.5 MB

```

```
In [40]: X.head()
```

```
Out[40]:
```

	Marital_Status	Gender_F	Gender_M	Age_0-17	Age_18-25	Age_26-35	Age_36-45	Age_46-50	Age_51-55	Age_55+	...	Pr
0	0	1	0	1	0	0	0	0	0	0	...	
1	0	1	0	1	0	0	0	0	0	0	...	
2	0	1	0	1	0	0	0	0	0	0	...	
3	0	1	0	1	0	0	0	0	0	0	...	
4	0	0	1	0	0	0	0	0	0	1	...	

5 rows × 76 columns

```
In [ ]:
```

```
In [41]: from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.30,random_state=42)
```

```
In [42]: print(X.shape)
print(y.shape)
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```

(547391, 76)
(547391,)
(383173, 76)
(383173,)
(164218, 76)
(164218,)

```

```
In [43]: from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)
```

```
In [44]: print(X.shape)
print(y.shape)
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```

(547391, 76)
(547391,)
(383173, 76)
(383173,)
(164218, 76)
(164218,)

```

```
In [45]: # Training Model selection
```

```
In [46]: from sklearn.linear_model import LinearRegression
LR=LinearRegression()
```

```
In [47]: LR.fit(X_train,y_train)
```

```
Out[47]: LinearRegression()
```

```
In [48]: y_pred=LR.predict(X_test)
```

```
In [49]: y_pred
```

```
Out[49]: array([11073.22265906, 13913.72265906, 7403.22265906, ...,
              7768.22265906, 5842.22265906, 7215.22265906])
```

```
In [50]: from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
In [51]: R2_score = r2_score(y_test,y_pred)
```

```
In [52]: print("training score = ",LR.score(X_train, y_train))
print("Testing score = ",LR.score(X_test, y_test))
print("Mean Absolute error =",mean_absolute_error(y_test,y_pred))
print("Mean Squared error =",mean_squared_error(y_test,y_pred))
print("Root Mean Squared error=",np.sqrt(mean_squared_error(y_test,y_pred)))
print("R2score = " ,R2_score)
```

```
training score = 0.6347153750403496
Testing score = 0.6350016170765882
Mean Absolute error = 2260.2782310954826
Mean Squared error = 8934811.918436665
Root Mean Squared error= 2989.1155746201357
R2score = 0.6350016170765882
```

```
In [53]: from sklearn.linear_model import Lasso
lasso = Lasso(alpha =0.0001)
lasso.fit(X_train, y_train)
y_pred = lasso.predict(X_test)
```

```
In [54]: print("training score =", lasso.score(X_train, y_train))
print("Testing score =", lasso.score(X_test, y_test))
print("Mean Absolute error =",mean_absolute_error(y_test,y_pred))
print("Mean Squared error =",mean_squared_error(y_test,y_pred))
print("Root Mean Squared error=",np.sqrt(mean_squared_error(y_test,y_pred)))
print("R2score = " ,R2_score)
```

```
training score = 0.6347213760940762
Testing score = 0.6350018580884018
Mean Absolute error = 2260.399604181418
Mean Squared error = 8934806.018697584
Root Mean Squared error= 2989.1145877496206
R2score = 0.6350016170765882
```

```
In [55]: from sklearn.linear_model import Ridge
Ridge = Ridge(alpha = 0.01)
Ridge.fit(X_train,y_train)
y_pred = Ridge.predict(X_test)
```

```
In [56]: print("training score =", Ridge.score(X_train, y_train))
print("Testing score =", Ridge.score(X_test, y_test))
print("Mean Absolute error =",mean_absolute_error(y_test,y_pred))
print("Mean Squared error =",mean_squared_error(y_test,y_pred))
print("Root Mean Squared error=",np.sqrt(mean_squared_error(y_test,y_pred)))
print("R2score = " ,R2_score)
```

```
training score = 0.6347213760941086
Testing score = 0.635001857336509
Mean Absolute error = 2260.3995908922634
Mean Squared error = 8934806.037103202
Root Mean Squared error= 2989.1145908283947
R2score = 0.6350016170765882
```

```
In [57]: from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor()
rf.fit(X_train,y_train)
```

```
Out[57]: RandomForestRegressor()
```

```
In [58]: print("training score =", rf.score(X_train, y_train))
print("Testing score =", rf.score(X_test, y_test))
print("Mean Absolute error =",mean_absolute_error(y_test,y_pred))
print("Mean Squared error =",mean_squared_error(y_test,y_pred))
print("Root Mean Squared error=",np.sqrt(mean_squared_error(y_test,y_pred)))
print("R2score = " ,R2_score)
```

```
training score = 0.7387019904768657
Testing score = 0.6314162482497752
Mean Absolute error = 2260.3995908922634
Mean Squared error = 8934806.037103202
Root Mean Squared error= 2989.1145908283947
R2score = 0.6350016170765882
```